
IF-Font: Ideographic Description Sequence-Following Font Generation

Xinping Chen, Xiao Ke*, Wenzhong Guo
Fujian Provincial Key Laboratory of Networking Computing
and Intelligent Information Processing,
College of Computer and Data Science, Fuzhou University, Fuzhou 350116, China
Engineering Research Center of Big Data Intelligence,
Ministry of Education, Fuzhou 350116, China
{221027017, kex, guowenzhong}@fzu.edu.cn

Abstract

Few-shot font generation (FFG) aims to learn the target style from a limited number of reference glyphs and generate the remaining glyphs in the target font. Previous works focus on disentangling the content and style features of glyphs, combining the content features of the source glyph with the style features of the reference glyph to generate new glyphs. However, the disentanglement is challenging due to the complexity of glyphs, often resulting in glyphs that are influenced by the style of the source glyph and prone to artifacts. We propose IF-Font, a novel paradigm which incorporates Ideographic Description Sequence (IDS) instead of the source glyph to control the semantics of generated glyphs. To achieve this, we quantize the reference glyphs into tokens, and model the token distribution of target glyphs using corresponding IDS and reference tokens. The proposed method excels in synthesizing glyphs with neat and correct strokes, and enables the creation of new glyphs based on provided IDS. Extensive experiments demonstrate that our method greatly outperforms state-of-the-art methods in both one-shot and few-shot settings, particularly when the target styles differ significantly from the training font styles. The code is available at <https://github.com/Stareven233/IF-Font>.

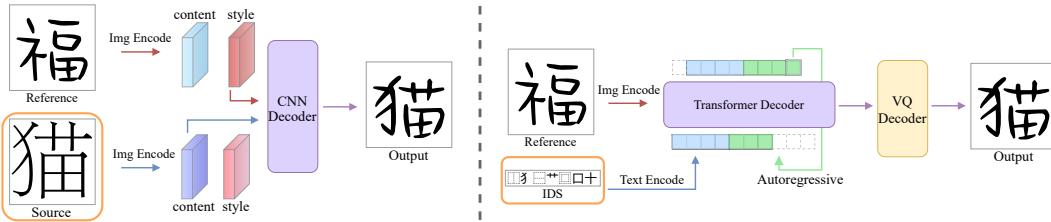


Figure 1: Comparison of two font generation paradigms. **Left:** The style-content disentangling paradigm. It assumes that a glyph can be decomposed into two distinct attributes: content and style. **Right:** The proposed paradigm. We first autoregressively predict the target tokens and decode them with a VQ decoder. Orange boxes show the main difference between the two paradigms.

*Corresponding author

1 Introduction

At the heart of font generation lies the extraction of styles from some reference glyphs of a certain font, and generate the remaining glyphs of this font. Some languages, such as Chinese, Japanese, and Korean, have a large number of characters and intricate glyph structures. Font generation can significantly reduce the labor intensity of font designers and support tasks like handwriting imitation, ancient book restoration, and internationalization of film and television productions.

EMD [58] and SA-VAE [44] are based on the belief that the target glyph can be generated by integrating the content features of the source glyph with the style features of the reference glyph, as illustrated in Fig. 1 (left). The majority of subsequent works [53, 38, 47, 25, 50] continues this paradigm, but this makes font generation a sub-task of image-to-image translation, where the source glyph is morphed to match the style of the reference glyphs, rather than being truly “generated”. Due to the complex structure of glyphs, achieving a distinct boundary between style and content features requires considerable effort. Consequently, glyphs produced through the disentangling strategy typically maintain similar stroke thickness to the reference glyph but align more closely with the content font regarding spatial structure, size, and inclination.

To this end, DG-Font [54] incorporates deformable convolution. Diff-Font [15] integrate diffusion process to improve the network’s learning capabilities. CF-Font [50] proposes content fusion. Additionally, several approaches [38, 47, 25, 30] combine fine-grained prior information such as strokes and components, to further enhance generation quality. These methods essentially follow the content-style disentanglement paradigm, in scenarios where the content font differs substantially from the target font, the resulting glyphs are susceptible to artifacts such as missing strokes, blur, and smudge.

We abandon the source glyphs in favor of Ideographic Description Sequence (IDS) to convey content information. It is based on a simple fact: without disentangling features, there is no risk of incomplete disentanglement. Consequently, font generation is reframed as sequence prediction task, where the objective is to generate the tokens of the target glyph based on the given IDS and reference glyphs. This approach mitigates the impact of source glyphs on the outputs and diminishes artifacts by leveraging the prior knowledge embedded in the quantized codebook. The users are allowed to formulate IDSs to create non-existing Chinese characters, such as *kokujī*² (Japanese-invented Chinese characters), provided that the corresponding structures and components have been learned during training. This endows the model with certain cross-linguistic capabilities. We refer to this method as **Ideographic Description Sequence-Following Font Generation**, or IF-Font. In summary, the key contributions of this paper are as follows:

- We propose IF-Font, which abandons the previous content-style disentanglement paradigm and generates glyphs through next-token prediction.
- We devise a novel IDS Hierarchical Analysis (IHA) module that analyzes the spatial structures and components of Chinese characters. It allows our decoder flexibly control the generated content with the encoded semantic features.
- Leveraging corresponding IDSs, we design the Structure-Style Aggregation (SSA) block to extract and efficiently aggregate the style features of reference glyphs.

2 Related Works

Image-to-image translation Image-to-image translation (I2I) aims to learn a mapping from a source domain to a target domain, requiring the transformation of images in the source domain into those in the reference style’s target domain while preserving their content. Pix2Pix [21] is the first I2I method that trains GANs [13] using paired data. CycleGAN [60] achieves unsupervised training through cycle consistency loss, although it is limited to transformations between two domains. UNIT [29] enforces the latent codes of images from two distinct domains to be identical, while employing separate generators for images in each domain. This process embodies the concept of disentanglement. MUNIT [19] further refines UNIT’s latent code into content and style codes. Multimodal image translation can be achieved by combining the content code with different style codes.

²<https://www.sljfaq.org/afaq/kokujii-list.html>

While applying an image-to-image translation framework for font generation is currently the mainstream approach, we believe this to be inappropriate. Unlike ordinary images, the boundaries between content and style in glyphs are ambiguous. For example, although handwriting will certainly differ when the same characters are written by different writers, their semantic meanings remain unchanged. Given that glyph features are challenging to disentangle, we utilize style-neutral IDSs to determine characters, thus avoiding any influence on the styles of results due to insufficient disentanglement of content glyphs.

Few-shot font generation Few-shot learning [12] represents the prevailing research focus in font generation, aiming to simulate the target style with just a handful of reference glyphs. Font generation methods fall into two categories based on their utilization of implicit structural information within glyphs. Methods treating glyphs as general images possess broader applicability, enabling generation across various languages. Conversely, methods leveraging structural information typically yield higher quality outputs but are confined to specific language.

Among the methods that do not incorporate structural information, EMD [58] stands out as the earliest attempt to disentangle glyphs into content and style features. DG-Font [54] employs deformable convolution [7] to capture the glyph deformations. FontRL [31] uses reinforcement learning [45] to draw the skeleton of Chinese characters. FontDiffuser [55] models the font generation task as a noise-to-denoise paradigm. Shamir et al. [42] explores a parametric representation of oriental alphabets, which can elegantly balance glyph quality and compression. In vector font generation, Deepvecfont [51] exhaustively exploit the dual-modality information (raster images and draw-command sequences) of fonts to synthesize vector glyphs. CLIPFont [43] controls the desired font style through text description rather than relying on style reference images.

In terms of methods that incorporate structural information, SA-VAE [44] utilizes the radicals and spatial structures of Chinese characters. CalliGAN [53] adopts the Zi2Zi framework [48] and fully decompose Chinese characters into sequences of components. SC-Font [23] further decomposes Chinese characters into stroke granularity. DM-Font [4] proposes dual memory to update component features continuously. LF-Font [38] represents component-wise style through low-rank matrix factorization [3]. MX-Font [37] automatically extracts the features through multiple localized experts. FS-Font [47] demands that reference glyphs include all components of the target glyph, otherwise may result in a degradation of generation quality. CG-GAN [25] employs GRU [6] and attention mechanism to predict component sequences. XMP-Font [30] performs multimodal pre-training on Chinese character strokes and glyphs data. Most of the above methods are constrained by the content-style disentanglement paradigm. They often neglect the presence of Ideographic Description Character (IDC) which refers to the spatial structure of Chinese characters, suffering from artifacts and inconsistent styles.

Vector quantized generative models Vector Quantization (VQ) typically follows a two-stage training scheme. Initially, it employs a codebook to record and update vectors, converting them from a continuous feature space to a discrete latent space. Subsequently, it models the distribution of these quantized vectors with a decoder to predict tokens, which are the codebook indices, and then restores the tokens to a image.

VQ-VAE [35] is the first to incorporate quantization into the variational autoencoder (VAE) [24] framework. VQ-VAE2 [41] performs multi-scale quantization and adopts rejection sampling [1] VQGAN [9] acquires the codebook with the help of GAN [13] and employs Transformer [49] to replace the PixelCNN [34] used by VQ-VAE [35]. RQ-VAE [26] proposes a residual quantizer. BEiT [2] performs masked image modeling (MIM) on the patch view with the supervision of visual tokens. MaskGIT [5] directly models visual tokens and proposes parallel decoding. MAGE [27] is similar to MaskGIT [5], but introduces ViT [8] and contrastive learning [14]. DQ-VAE [18] further encodes images with variable-length codes.

Since quantization is equivalent to tokenizing images, many methods attempt to enable multimodal generation. The dVAE proposed by DALL-E [40] relaxes the discrete sampling problem utilizing Gumbel-Softmax [33, 22], outputing the probability distribution of codebook codes. SEED [11] designed a Causal Q-Former to extract image embeddings and quantize them. LQAE [28] trains VQ-VAE [35] to quantize the image into the frozen LLM codebook space directly. SPAE [56] introduces multi-layer and coarse-to-fine pyramid quantization and semantic guidance with CLIP

[39]. V2L [61] further proposes global and local quantizers. Given the absence of a pre-trained model tailored for IDS, we directly concatenate visual tokens with IDS tokens to performed autoregression.

3 Method

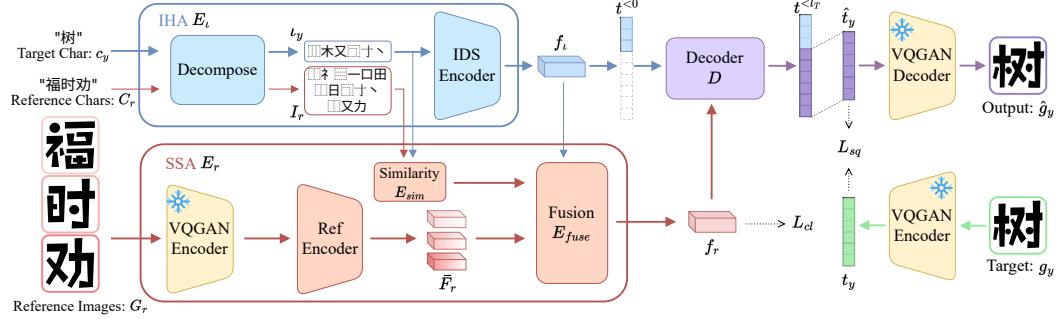


Figure 2: Overview of our proposed method. The overall framework mainly consists of three parts: IDS Hierarchical Analysis module E_t , Structure-Style Aggregation block E_r , and a decoder D .

As shown in Fig. 2, given a target character c_y , reference characters $C_r = \{c_r^i\}_{i=1}^k$ and reference glyphs $G_r = \{g_r^i\}_{i=1}^k$, the goal of our framework is to generate a glyph \hat{g}_y that conforms to the semantics of c_y and has a style consistent with G_r . To achieve this objective, we analyze c_y with IHA to derive its associated IDS ι_y , which is then encoded as a semantic feature f_t . Likewise, we can obtain the IDS $I_r = \{\iota_r^i\}_{i=1}^k$ corresponding to C_r . Following this, we employ the similarity module E_{sim} to assess the relationship between ι_y and I_r . Combined with f_t and the output of E_{sim} , the features \bar{F}_r corresponding to G_r are fused into the final style feature f_r in the SSA block. ι_y is reshaped as initial tokens $t^{<0}$, which is fed into the decoder D along with f_r for autoregressive modeling. Finally, the predicted glyph tokens \hat{t}_y are decoded with the pre-trained VQGAN to obtain the generated glyph \hat{g}_y .

3.1 IDS Hierarchical Analysis

A simple alternative to using a source glyph as input is to directly employ the character itself to control the semantics of the output. However, considering the vast number of characters in Chinese, this approach proves to be impractical due to its expensive cost. Moreover, it overlooks the structural intricacies of characters.

Ideograph Description Sequence is a structural description grammar for Chinese characters defined by the Unicode Standard, which consists of description characters and basic components (mainly Chinese characters) through a prefix notation. Decomposing Chinese characters into their corresponding IDSs can notably streamline the vocabulary, allowing characters with similar structures or components to share common features.

However, a Chinese character may have multiple equivalent IDSs. Many Chinese characters have a top-bottom or left-right structure, the IDCs follows a long-tail distribution, presenting challenges for model training. Fortunately, the left-middle-right structure of Chinese characters can be equivalently represented by two left-right structures. Similarly, the top-middle-bottom structure equals two top-bottom structures. The examples can be found in Fig. 3.

Based on the above observation, we employ a IDS Hierarchical Analysis (IHA) module. Instead of rigidly querying the decomposition table when determining the IDS of a character, we examine whether the character follows a left-middle-right or top-middle-bottom structure. Subsequently, we construct multiple equivalent IDSs for the same character through random selection. To summarize, c_y and C_r are initially decomposed into ι_y and I_r , respectively. In the IDS encoder, ι_y is padded to the maximum sequence length l_I and encoded into the associated semantic feature $f_t \in \mathbb{R}^{l \times c}$.

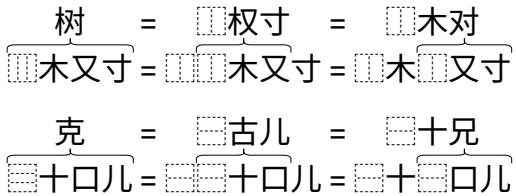


Figure 3: The illustration of equivalent IDSs.

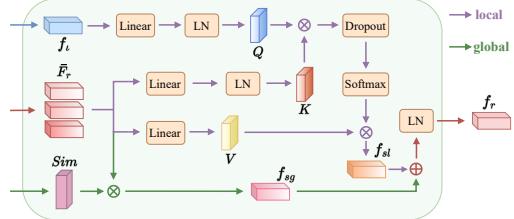


Figure 4: The illustration of the fusion module E_{fuse} of SSA block.

3.2 Structure-Style Aggregation

Many previous methods [55, 10, 59, 25, 30, 54] overlook interactions between reference and target characters during the extraction of reference styles, resulting in a lack of relevance in the extracted features. The more closely the reference character resembles the target character, the more effortlessly the generation process can preserve the style. Ideally, employing the target glyph itself as the reference, known as self-reconstruction, should yield the most effective output. Although FS-Font [47] endeavors to ensure that the reference characters cover all components of the target character, its implementation hinges on predefined content-reference mapping, which may limit its adaptability. To address this issue, we propose a Structure-Style Aggregation (SSA) block, as shown in Fig. 2. We convert the reference glyph G_r into the latent space of VQGAN and encode it one by one into the corresponding intermediate features $\bar{F}_r = \{\bar{f}_r^i \in \mathbb{R}^{h \cdot w \times c}\}_{i=1}^k$. The similarity module E_{sim} evaluates the resemblance between each reference IDS I_r and the target IDS ι_y , considering whether they share identical structures or components. It produces a set of similarity weights $Sim = \{sim^i \in \mathbb{R}^1\}_{i=1}^k$, which can guide the subsequent feature fusion process. The fusion module E_{fuse} consists of two branches: global and local style feature aggregation, as shown in Fig. 4. The global features mainly focus on the glyph layout, stroke thickness, and inclination, which can be obtained by merging the coarse style features \bar{F}_r with the similarity weight Sim obtained in the previous step:

$$f_{rg} = \text{softmax}(Sim) \bar{F}_r \in \mathbb{R}^{h \cdot w \times c}. \quad (1)$$

While local features are more concerned with the strokes, such as stroke length, stroke edge, and other nuances, we adopt cross-attention to gather the required style feature according to the needs of the target character:

$$\begin{aligned} F'_s &= \text{flatten}^2(\bar{F}_r) \in \mathbb{R}^{k \cdot h \cdot w \times c}, \quad Q = \text{LayerNorm}(L_q(f_r)) \in \mathbb{R}^{l \times c}, \\ K &= \text{LayerNorm}(L_k(F'_s)) \in \mathbb{R}^{k \cdot h \cdot w \times c}, \quad V = L_v(F'_s) \in \mathbb{R}^{k \cdot h \cdot w \times c}, \\ A &= \text{dropout}\left(\frac{QK^T}{\sqrt{c}}\right) \in \mathbb{R}^{l \times k \cdot h \cdot w}, \quad f_{rl} = \text{softmax}(A)V \in \mathbb{R}^{l \times c}, \end{aligned} \quad (2)$$

where $\text{flatten}^2(\cdot)$ denotes flattening the first two dimensions of the feature, and L_q , L_k , L_v are linear projections, and $\text{LayerNorm}(\cdot)$ denotes layer normalization. In Eq. 3, we obtain the aggregated style feature based on Eq. 1 and Eq. 2, where \circ denotes concatenation operator:

$$f_r = \text{LayerNorm}(f_{rg} \circ f_{rl}). \quad (3)$$

3.3 Style Contrast Enhancement

There are some strategies to maintain style consistency: integrating consistency loss [59, 25], introducing a discriminator to determine the generated style [25, 47, 36], or treating the extracted style feature as a variable for further optimization [50]. These approaches are indeed beneficial for improving the generation quality, but they may be inflexible or introduce additional parameters.

In this paper, we propose a streamlined approach named the Style Contrast Enhancement (SCE) module, which promotes the proximity of representations for the same style and the distance between representations for different styles. We apply a linear projection to the style feature f_r , resulting in a contrastive feature $e = MLP(f_r)$.

In one batch, we denote the indices of contrastive features corresponding to all samples as $E_* = \{i \in \mathbb{N} \mid 0 \leq i < 2N\}$, where N represents the batch size. The dimensionality of E_* is double the

batch size N due to our utilization of a momentum encoder [16]. Each sample x_a within the batch undergoes processing by both the encoder and the momentum encoder, yielding two outputs that serve as positive pairs. The negative sample set is defined as $E_- = \{i \in E_* \mid s(x_i) \neq s(x_a)\}$, while the positive sample set is $E_+ = \{i \in E_* \mid i \neq a, s(x_i) = s(x_a)\}$, where $s(\cdot)$ denotes the operator used to retrieve the corresponding style. The contrastive loss can be calculated as follows:

$$\mathcal{L}_{cl} = -\frac{1}{2N} \sum_{a \in E_*} \log \frac{\sum_{p \in E_+} \exp(e_a^T e_p / \tau)}{\sum_{p \in E_+} \exp(e_a^T e_p / \tau) + \sum_{n \in E_-} \exp(e_a^T e_n / \tau)}. \quad (4)$$

3.4 Generation

The decoder D is provided with both semantic feature f_s and style feature f_r . It treats f_s as the initial tokens $t^{<0} = f_s$, and then predicts the distribution of the next token autoregressively as $p(t^i \mid t^{<i}, f_r)$. Each newly predicted token is appended to the previous tokens for the subsequent iteration, it continues until all tokens are predicted. The likelihood of the entire sequence can be calculated as $\prod_{i=0}^{l_T-1} p(t^i \mid t^{<i}, f_r)$. There are two ways for incorporating f_r . The most straightforward approach involves using f_r as initial tokens, represented by $t^{<0} = f_r \circ f_s$, akin to f_s . These tokens participate in each forward pass, relying on the self-attention mechanism to extract and integrate features.

However, it is only practical for low-resolution scenarios. Viewing f_r as tokens may lead to excessively long sequences, requiring a balance between computational efficiency and generation quality. To address this challenge, we incorporate f_r into each block of the decoder D via cross-attention. The tokens act as queries to align with the corresponding style features. $t^{<l_T}$ denotes all the predicted tokens, from which $t^{<0}$ is removed to get the glyph tokens \hat{t}_y . The objective in Eq. 5 is to maximize the log-likelihood of the token sequence.

$$\mathcal{L}_{sq} = -\log \left(\prod_{i=0}^{l_T-1} p(t^i \mid t^{<i}, f_r) \right), \quad (5)$$

Finally, the model can be trained according to the objective in Eq. 6.

$$\mathcal{L}_{total} = \mathcal{L}_{sq} + \lambda_{cl} \mathcal{L}_{cl}, \quad (6)$$

where λ_{cl} controls the weight of contrastive loss, cf. Eq. 4, we set $\lambda_{cl} = 0.5$ in our experiments.

4 Experiments

4.1 Dataset and Evaluation Metrics

Datasets We gathered 464 fonts from the Internet, covering diverse categories like printed, handwritten, and artistic styles. Next, we selected 3,500 commonly encountered Chinese characters and rendered them into 128x128 resolution images using the collected fonts.

The training set comprises 3,300 randomly selected Chinese characters and 424 fonts, referred to as Seen Fonts and Seen Characters (**SFSC**). There are two test sets: the first includes the same 3,300 characters but with different 40 fonts, called Unseen Fonts and Seen Characters (**UFSC**). The second test set consists of the remaining 200 characters and the same 40 fonts, known as Unseen Fonts and Unseen Characters (**UFUC**). We found a publicly accessible IDS decomposition table³. However, it exhibits several redundant entries and circular references, as well as an absence of some characters. Therefore, we performed simplifications and enhancements, reducing the number of IDCs to the 12 depicted in Table 1, which is sufficient for most frequently used Chinese characters. For convenience, we set the basic component’s IDS as itself.

Table 1: 12 IDCs used in this paper.

IDC	Structure	Example (Char:IDS)
□	left-right	鸿:□江鸟, 鸣:□𠂇丶
■	top-bottom	惹:□若心, 主:□丶王
□■	left-middle-right	鸿:□沙工鸟, 小:□𠂇丶
■□	top-middle-bottom	惹:□𠂇右心, 叁:□ム大三
□□	enclosed-surrounding	回:□口口, 又:□又丶
□■■	left-top-right-surrounding	闪:□门人, 太:□大丶
□■□	left-bottom-right-surrounding	山:□口丨, 义:□义丶
□■□■	top-left-bottom-surrounding	匠:□匚斤, 兔:^□免丶
□■□■	top-left-surrounding	友:□大又, 厌:□厂犬
□■□■	top-right-surrounding	乃:□乃ノ, 勹:□匚丶
□■□■	left-bottom-surrounding	边:□辵力, 犬:□大丶
□■□■	overlaying	平:□干丶, 丸:□九丶

³<https://babelstone.co.uk/CJK/IDS.TXT>

Evaluation metrics We compare all methods in the following metrics, i.e., FID [17], L1, LPIPS [57], RMSE, and SSIM [52]. Since aesthetics is inherently subjective, we conduct a user study for all methods to evaluate model performance based on user satisfaction. We observe that the existing font generation methods have differences in data preprocessing and metric selection. Factors such as glyph resolution, the padding around glyphs, the range of pixel values, the number of reference glyphs, and the evaluation function implementation all influence metric values. For example, NTF [10] and CF-Font [50] center the glyph within the canvas, leaving white space around it. However, this leads to inflated metric calculations. To ensure a fair comparison, we adopt consistent test data and metric implementation across all methods under evaluation. Specifically, we eliminate padding around the glyphs, fix the canvas resolution to 128 pixels, scale the data range to $[0, 1]$, utilize SqueezeNet [20] as the network type to calculate LPIPS [57], and select the inceptionv3 [46] feature layer with 2048 dimensions for FID [17] calculation.

Table 2: Quantitative evaluation on UFSC and UFUC dataset. “User” indicates user study, the samples are generated under 3-shot setting. Bold and underlined numbers denote the best and the second best respectively. Please refer to Fig. 10 in Appendix for the corresponding radar plots.

Methods	1shot					3shot					8shot					User (%) \uparrow	
	FID \downarrow	L1 \downarrow	LPIPS \downarrow	RMSE \downarrow	SSIM \uparrow	FID \downarrow	L1 \downarrow	LPIPS \downarrow	RMSE \downarrow	SSIM \uparrow	FID \downarrow	L1 \downarrow	LPIPS \downarrow	RMSE \downarrow	SSIM \uparrow		
UFSC	CG-GAN [25]	11.3911	0.1784	0.1500	0.3997	0.4428	10.8713	0.1771	0.1464	0.3982	0.4441	11.1332	0.1764	0.1457	0.3974	0.4440	14.78
	LF-Font [38]	32.9264	0.1764	0.1586	0.3967	0.4465	29.1840	0.1786	0.1576	0.3998	0.4432	26.9590	0.1694	0.1567	0.3875	0.4590	11.42
	FS-Font [47]	112.0971	0.2836	0.3108	0.5145	0.2795	25.5231	0.2075	0.1916	0.4343	0.3865	93.7912	0.1900	0.2086	0.4124	0.4183	6.55
	CF-Font [50]	20.4457	0.1839	0.1581	0.4066	0.4323	30.8426	0.1767	0.1650	0.3977	0.4468	30.9829	0.1784	0.1595	0.3990	0.4465	12.22
	VQ-Font [36]	72.7064	0.1958	0.2215	0.4201	0.4077	32.9390	0.1789	0.1775	0.4016	0.4405	33.6378	0.1774	0.1732	0.3995	0.4413	7.50
	NTF [10]	35.3797	0.2602	0.2027	0.4887	0.3244	26.1215	0.2275	0.1749	0.4542	0.3659	23.0519	0.2238	0.1739	0.4501	0.3720	6.43
	FontDiffuser [55]	3.9969	0.1938	0.1371	0.4180	0.4076	3.6989	0.1774	0.1248	0.3980	0.4370	4.1017	0.1748	0.1234	0.3947	0.4420	18.32
	IF-Font (Ours)	6.7695	0.1520	<u>0.1307</u>	0.3688	0.4915	6.8359	0.1478	0.1258	0.3620	<u>0.5021</u>	6.7383	0.1429	0.1216	0.3552	<u>0.5140</u>	22.78
UFUC	CG-GAN [25]	13.4734	0.1805	0.1508	0.4019	0.4362	13.0347	0.1790	0.1471	0.4001	0.4383	13.2049	0.1780	0.1462	0.3991	0.4391	15.73
	LF-Font [38]	37.3840	0.1835	0.1620	0.4047	0.4283	28.8252	0.1850	0.1609	0.4071	0.4283	30.5147	0.1735	0.1582	0.3920	0.4473	11.65
	FS-Font [47]	112.6636	0.2847	0.3112	0.5155	0.2764	31.2833	0.2106	0.1923	0.4373	0.3785	98.9484	0.1921	0.2095	0.4146	0.4131	6.17
	CF-Font [50]	22.8601	0.1865	0.1584	0.4094	0.4259	34.0245	0.1796	0.1660	0.4009	0.4399	33.2477	0.1809	0.1601	0.4019	0.4399	12.03
	VQ-Font [36]	75.1737	0.1980	0.2217	0.4223	0.4018	36.4831	0.1809	0.1776	0.4037	0.4345	36.5486	0.1796	0.1733	0.4017	0.4354	8.23
	NTF [10]	39.3581	0.2678	0.2074	0.4958	0.3086	29.9205	0.2303	0.1753	0.4568	0.3593	27.9580	0.2290	0.1755	0.4553	0.3619	6.33
	FontDiffuser [55]	8.2524	0.1914	0.1527	0.4157	0.4163	7.6444	0.1771	0.1413	0.3981	0.4418	8.9166	0.1702	0.1367	0.3890	0.4543	18.57
	IF-Font (Ours)	8.4844	0.1651	<u>0.1387</u>	0.3845	0.4676	8.4922	<u>0.1597</u>	0.1338	0.3775	0.4782	8.3203	0.1561	0.1305	0.3728	0.4864	21.28

4.2 Comparison with state-of-the-art Methods

We compare the proposed IF-Font with seven SOTA methods on our UFSC and UFUC datasets respectively, including CG-GAN [25](CVPR 2022), LF-Font [38](TPAMI 2022), FS-Font [47](CVPR 2022), CF-Font [50](CVPR 2023), VQ-Font [36](ICCV 2023), NTF [10](CVPR 2023) and FontDiffuser [55] (AAAI 2024). All methods are trained from scratch on our SFSC dataset according to their respective official codes and default configurations. We slightly modify the codes of CG-GAN, LF-Font, FS-Font, VQ-Font and FontDiffuser to support varied numbers of reference glyphs.

4.2.1 Quantitative comparison

Table 2 compares IF-Font and other SOTA methods. IF-Font significantly surpasses competitors in all reference glyph number settings for both datasets. Notably, IF-Font’s performance on FID metric is exceptionally low, reaching a single-digit score, thanks to the high quality and clarity of the samples it generates. FS-Font [47] relies heavily on the predefined content-reference mapping, whereas the reference glyphs in all our experiments are randomly selected. Especially, when only one reference glyph is provided, covering all components of the target character becomes challenging, leading to poor performance of FS-Font, as shown in Fig. 10a. NTF [10] also struggles to imitate the target style, the layout of its generated samples often resembles that of the source font. In cases where there’s a significant disparity between the source and target styles, NTF is prone to missing strokes.

We attribute CF-Font’s performance to its reliance on fusing contents of 10 basic fonts. However, there happen to be a gap between the train dataset and our evaluation dataset. We conduct a user study through Fuxi Youling Crowdsourcing Platform ⁴ to quantify the subjective quality. For each test dataset, 5 characters are randomly selected, and each model is required to generate glyphs corresponding to 40 unseen fonts. A total of 30 participants are asked to select the option that most closely matches the ground truth from the generated results. The outcomes of the user study are presented in the last column of Table 2.

⁴<https://fuxi.163.com/solution/data>

Figure 5: Qualitative comparison with state-of-the-art methods, in which red boxes outline the artifacts. “Source” denotes the content glyph of other methods, IF-Font only employs the corresponding IDS.

4.2.2 Qualitative comparison

We present the corresponding samples from Table 2 in Fig. 5. IF-Font stands out by producing the clearest and most style-consistent samples. In contrast, FS-Font [47], LF-Font [38], CF-Font [50], and other models exhibit issues such as stroke errors or blur. VQ-Font [36] and NTF [10] are constrained by the source font and struggle with flat or narrow layouts, resulting in incorrect structures. VQ-Font even tends to crop marginal parts of glyphs to fit the target style. While CF-Font generally preserves the correct glyph layout, its outputs exhibit noticeable artifacts, indicating some remaining style inconsistencies. The performance of FontDiffuser [55] is also outstanding, but there is still a slight deficiency in the imitation of font styles. On the other hand, IF-Font maintains the correct character structures and excels in aspects such as the aspect ratio, glyph layouts, and stroke details.

Table 3: Ablation studies on different modules. The first row is the results of baseline. I, S and C represent IHA, SSA, and SCE respectively.

Module	FID \downarrow	L1 \downarrow	LPIPS \downarrow	RMSE \downarrow	SSIM \uparrow		
I	S	C					
\times	\times	\times	8.2656	0.1632	0.1383	0.3820	0.4728
\checkmark	\times	\times	8.3750	0.1638	0.1381	0.3828	0.4764
\checkmark	\checkmark	\times	7.5391	0.1614	0.1348	0.3797	0.4780
\checkmark	\checkmark	\checkmark	8.4922	0.1597	0.1338	0.3775	0.4782

IDS	Reference	Baseline	+I	+IS	+ISC	Target
四四西示几×	福	飄	飄	飄	飄	飄
四木木彑	福	移	彬	彬	彬	彬
四ム大三	福	參	參	參	參	參
四氵一丁一鳥	福	鴻	鴻	鴻	鴻	鴻

Figure 6: Visualization of different modules in Table 3. Red, blue and green boxes represent the missing components, style inconsistency and corresponding improvements respectively.

4.3 Ablation Studies

Main modules Removing the IHA, SSA, and SCE modules of IF-Font, a baseline model can be obtained. For a input character, it directly looks up the decomposition table to derive the corresponding IDS, and then encode the semantic feature f_t through a embedding layer. The intermediate style features \bar{F}_r are directly averaged as style features f_r , excluding any interactions with the similarity weight Sim and semantic feature f_t . The whole model relies solely on cross-entropy loss for supervision. Building upon the baseline, we incrementally reintegrate three modules to assess their individual contribution. Quantitative results are presented in Table 3, while Fig. 6 provides visualizations of these results. For further ablation study of the SSA block, please refer to Table 6 in Appendix B.1. Upon integrating our modules, a consistent improvement is observed across most metrics. Fig. 6 illustrates how IHA alleviates the issue of missing components present in the baseline. SSA enhances style consistency, while SCE improves the capability to imitate styles.

Table 4: The impact of IDS granularity on performance.

Granularity	FID↓	L1↓	LPIPS↓	RMSE↓	SSIM↑
Component	8.4922	0.1597	0.1338	0.3775	0.4782
Stroke	8.4297	0.1616	0.1347	0.3799	0.4888
Mixed	8.5234	0.1598	0.1337	0.3775	0.4782

IDS granularity We further analyze the impact of three different IDS granularities: components, strokes and mixed. Please see Table 5 in the Appendix for the examples of these granularities. Table 4 shows the quantitative results. Stroke granularity results in performance degradation across three metrics. We attribute this decline to the conflict between IDSs, hindering the model’s ability to identify the target character. An attempt to concatenate IDSs from both granularities yields performance akin to that of component granularity. While this approach extends the sequence considerably, hence we opt for component granularity.

4.4 Visualization of SSA



Figure 7: Visualization of attention maps between IDS and reference glyphs. The symbols above are the target character (black) and the corresponding IDS (orange).

To demonstrate the effectiveness of Structure-Style Aggregation block, we visualize the attention maps corresponding to each IDC and component within the IDSs relative to the reference glyphs, as depicted in Fig. 7. Specifically, we choose the matrix $A \in \mathbb{R}^{l \times k \cdot h \cdot w}$ in local feature calculation. For each position i of the target $IDS_{t,y}$, there exists a corresponding attention map $A^i \in \mathbb{R}^{k \cdot h \cdot w}$, which indicates the attention that t_y^i pays to the k reference features. We present the attention map A^i to visualize the distribution of attention weights directly. Additionally, we apply opacity to this map and overlay it onto the original reference glyph.

As we can see, when the target IDC or component exists in the reference glyph, more attention will be paid to the corresponding place. For instance, in the first row, the first, second, and fourth columns, and in the second row, the third and fifth columns are distinctly highlighted. Conversely, if the reference glyph lacks the target component, the local branch tends not to engage, as evidenced by the nearly blank third row. This approach stems from a preference to avoid forced attention allocation which might lead to interference. Instead, leveraging the average style captured by the global branch helps maintain a baseline quality of the output.

4.5 New Glyph Creation

We validate the flexibility of IF-Font by generating the glyphs using IDSs of non-standard Chinese characters. Fig. 8 shows our experimental results. IF-Font demonstrates robust generalization by following given IDSs to produce new glyphs with accurate structure and consistent style. We fixed the font to “Sarasa Gothic” in the experiment, which is a CJK programming font. The last two columns lack ground truths due to their entirely non-existent characters.

IDS	Output	Target	IDS	Output	Target	IDS	Output	IDS	Output
日雨下	雰	雰	日丶	丶丶	丶丶	日口色	色	日乃马	局
日火田	烟	烟	日艸世	世	世	日广祭	祭	日口木	困
日几巾	廴	廴	日入辶	辶	辶	日气杀	杀	日雨木心	丽
日米千	杆	杆	日白田	田	田	日是夏	夏	日雨木心	悉
日山𠂇	嶮	嶮	日口天	天	天	日口云	云	廿人	廿

Figure 8: The ability of the IF-Font to create glyphs. The first two columns are kokuj, and the last two columns are completely non-existent Chinese characters.

5 Discussion



Figure 9: Failed cases on complex fonts of UFUC. Orange boxes highlight reconstruction errors of VQGAN, red outlines the structural errors. *GT*: the glyphs rendered by fonts; *Target*: the glyphs reconstructed by VQGAN; *Output*: the glyphs generated by IF-Font.

Failure cases Although our method enables high-quality generation under most circumstances, it still struggles on some hard cases, as illustrated in Fig. 9. IF-Font encounters difficulties with fancy and irregular font styles, including those with decorations, extremely flat or narrow layouts, excessively curved strokes, and calligraphic writings. Despite these challenges, it continues to preserve the correct character structure. Further discussion on the reasons for the difficulties in generating these fonts can be found in Appendix C.3.

Usability We focus on CJK characters due to their unique spatial structures, which better reflect the characteristics of our method. By expanding the vocabulary and incorporating relevant data for training, IF-Font can also be adapted to handle other character sets.

Advantages *Conforms to writing habits.* We believe that the process of autoregressive modeling with IDS implicitly contains the order of writing. *Scalability.* Good scalability can be achieved by leveraging the mature experience of LLMs. *Robustness.* Due to vector quantization, glyphs are represented by a limited number of tokens (only 256 types), which reduces the learning difficulty for the decoder and decreases the likelihood of artifacts and other issues in the generated results.

6 Conclusion

We have presented IF-Font, a novel font generation paradigm. IF-Font redefines font generation as a sequence prediction task by quantizing glyphs as token sequences and leveraging Ideographic Description Sequence (IDS) to control the semantics of the generated glyphs. This method demonstrates exceptional capability in managing complex styles while preserving the correct structures. Benefiting from the flexibility of IDS, our method enables the creation of glyphs. This is achieved by formulating legal IDSs, which is a salient advantage over other methods that typically require the character to be present in at least one font as a precondition for generation. Refining and improving the IDS decomposition rules is considered future works. Furthermore, exploring the integration of IDS into handwritten font generation may yield interesting insights.

Acknowledgments and Disclosure of Funding

This work was supported in part by the National Natural Science Foundation of China under Grant 61972097 and U21A20472, in part by the National Key Research and Development Plan of China under Grant 2021YFB3600503, in part by the Natural Science Foundation of Fujian Province under Grant 2021J01612 and 2020J01494, in part by the Major Science and Technology Project of Fujian Province under Grant 2021HZ022007.

References

- [1] Samaneh Azadi, Catherine Olsson, Trevor Darrell, Ian J. Goodfellow, and Augustus Odena. Discriminator Rejection Sampling. In *International Conference on Learning Representations (ICLR)*, 2019.
- [2] Hangbo Bao, Li Dong, Songhao Piao, and Furu Wei. Beit: Bert Pre-Training of Image Transformers. In *International Conference on Learning Representations (ICLR)*, 2022.
- [3] Emmanuel J. Candès and Benjamin Recht. Exact Matrix Completion via Convex Optimization. *Foundations of Computational Mathematics*, 9(6):717–772, 2009.
- [4] Junbum Cha, Sanghyuk Chun, Gayoung Lee, Bado Lee, Seonghyeon Kim, and Hwalsuk Lee. Few-Shot Compositional Font Generation with Dual Memory. In *European Conference on Computer Vision (ECCV)*, pages 735–751, 2020.
- [5] Huiwen Chang, Han Zhang, Lu Jiang, Ce Liu, and William T. Freeman. Maskgit: Masked Generative Image Transformer. In *Computer Vision and Pattern Recognition (CVPR)*, pages 11305–11315, 2022.
- [6] Kyunghyun Cho, Bart van Merriënboer, Çağlar Gülcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1724–1734, 2014.
- [7] Jifeng Dai, Haozhi Qi, Yuwen Xiong, Yi Li, Guodong Zhang, Han Hu, and Yichen Wei. Deformable Convolutional Networks. In *IEEE International Conference on Computer Vision (ICCV)*, pages 764–773, 2017.
- [8] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. In *International Conference on Learning Representations (ICLR)*, 2021.
- [9] Patrick Esser, Robin Rombach, and Bjorn Ommer. Taming Transformers for High-Resolution Image Synthesis. In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 12873–12883, 2021.
- [10] Bin Fu, Junjun He, Jianjun Wang, and Yu Qiao. Neural Transformation Fields for Arbitrary-Styled Font Generation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [11] Yuying Ge, Yixiao Ge, Ziyun Zeng, Xintao Wang, and Ying Shan. Planting a SEED of Vision in Large Language Model. *ArXiv*, abs/2307.08041, 2023.
- [12] Abel Gonzalez-Garcia, Joost van de Weijer, and Yoshua Bengio. Image-to-image translation for cross-domain disentanglement. In *Conference on Neural Information Processing Systems (NeurIPS)*, pages 1294–1305, 2018.
- [13] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial networks. *Communications of the ACM*, 63(11):139–144, 2020.

- [14] Raia Hadsell, Sumit Chopra, and Yann LeCun. Dimensionality Reduction by Learning an Invariant Mapping. In *Computer Vision and Pattern Recognition (CVPR)*, volume 2, pages 1735–1742, 2006.
- [15] Haibin He, Xinyuan Chen, Chaoyue Wang, Juhua Liu, Bo Du, Dacheng Tao, and Y. Qiao. Diff-Font: Diffusion Model for Robust One-Shot Font Generation. *ArXiv*, abs/2212.05895, 2022.
- [16] Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, and Ross B. Girshick. Momentum Contrast for Unsupervised Visual Representation Learning. In *Computer Vision and Pattern Recognition (CVPR)*, pages 9726–9735, 2020.
- [17] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. Gans Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium. In *Conference on Neural Information Processing Systems (NeurIPS)*, pages 6626–6637, 2017.
- [18] Mengqi Huang, Zhendong Mao, Zhuowei Chen, and Yongdong Zhang. Towards Accurate Image Coding: Improved Autoregressive Image Generation with Dynamic Vector Quantization. In *Computer Vision and Pattern Recognition (CVPR)*, pages 22596–22605, 2023.
- [19] Xun Huang, Ming-Yu Liu, Serge J. Belongie, and Jan Kautz. Multimodal Unsupervised Image-to-Image Translation. In *European Conference on Computer Vision (ECCV)*, pages 179–196, 2018.
- [20] Forrest N. Iandola, Song Han, Matthew W. Moskewicz, Khalid Ashraf, William J. Dally, and Kurt Keutzer. SqueezeNet: Alexnet-level accuracy with 50x fewer parameters and <0.5mb model size. *arXiv*, 2016.
- [21] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A. Efros. Image-to-Image Translation with Conditional Adversarial Networks. In *Computer Vision and Pattern Recognition (CVPR)*, pages 5967–5976, 2017.
- [22] Eric Jang, Shixiang Gu, and Ben Poole. Categorical Reparameterization with Gumbel-Softmax. In *International Conference on Learning Representations (ICLR)*, 2017.
- [23] Yue Jiang, Zhouhui Lian, Yingmin Tang, and Jianguo Xiao. Scfont: Structure-Guided Chinese Font Generation via Deep Stacked Networks. *Proceedings of the AAAI Conference on Artificial Intelligence*, 33(01):4015–4022, 2019.
- [24] Diederik P. Kingma and Max Welling. Auto-Encoding Variational Bayes. In *International Conference on Learning Representations (ICLR)*, 2014.
- [25] Yuxin Kong, Canjie Luo, Weihong Ma, Qiyuan Zhu, Shenggao Zhu, Nicholas Yuan, and Lianwen Jin. Look Closer To Supervise Better: One-Shot Font Generation via Component-Based Discriminator. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 13482–13491, 2022.
- [26] Doyup Lee, Chiheon Kim, Saehoon Kim, Minsu Cho, and Wook-Shin Han. Autoregressive Image Generation using Residual Quantization. In *Computer Vision and Pattern Recognition (CVPR)*, pages 11513–11522, 2022.
- [27] Tianhong Li, Huiwen Chang, Shlok Mishra, Han Zhang, Dina Katabi, and Dilip Krishnan. Mage: Masked Generative Encoder to Unify Representation Learning and Image Synthesis. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [28] Hao Liu, Wilson Yan, and Pieter Abbeel. Language Quantized AutoEncoders: Towards Unsupervised Text-Image Alignment. In *Conference on Neural Information Processing Systems (NeurIPS)*, 2023.
- [29] Ming-Yu Liu, Thomas M. Breuel, and Jan Kautz. Unsupervised Image-to-Image Translation Networks. In *Conference on Neural Information Processing Systems (NeurIPS)*, pages 700–708, 2017.

- [30] Wei Liu, Fangyue Liu, Fei Ding, Qian He, and Zili Yi. Xmp-Font: Self-Supervised Cross-Modality Pre-Training for Few-Shot Font Generation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7905–7914, 2022.
- [31] Yitian Liu and Zhouhui Lian. Fontrl: Chinese Font Synthesis via Deep Reinforcement Learning. In *AAAI Conference on Artificial Intelligence (AAAI)*, pages 2198–2206, 2021.
- [32] Ilya Loshchilov and Frank Hutter. Decoupled Weight Decay Regularization. In *International Conference on Learning Representations (ICLR)*, 2019.
- [33] Chris J. Maddison, Andriy Mnih, and Yee Whye Teh. The Concrete Distribution: A Continuous Relaxation of Discrete Random Variables. In *International Conference on Learning Representations (ICLR)*, 2017.
- [34] Aaron van den Oord, Nal Kalchbrenner, and Koray Kavukcuoglu. Pixel Recurrent Neural Networks. *arXiv*, 2016.
- [35] Aäron van den Oord, Oriol Vinyals, and Koray Kavukcuoglu. Neural Discrete Representation Learning. In *Conference on Neural Information Processing Systems (NeurIPS)*, pages 6306–6315, 2017.
- [36] Wei Pan, Anna Zhu, Xinyu Zhou, Brian Kenji Iwana, and Shilin Li. Few Shot Font Generation Via Transferring Similarity Guided Global Style and Quantization Local Style. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, volume abs/2309.00827, pages 19506–19516, 2023.
- [37] Song Park, Sanghyuk Chun, Junbum Cha, Bado Lee, and Hyunjung Shim. Multiple Heads are Better than One: Few-shot Font Generation with Multiple Localized Experts. In *IEEE International Conference on Computer Vision (ICCV)*, pages 13880–13889, 2021.
- [38] Song Park, Sanghyuk Chun, Junbum Cha, Bado Lee, and Hyunjung Shim. Few-shot Font Generation with Weakly Supervised Localized Representations. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, PP:1–17, 2022.
- [39] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning Transferable Visual Models From Natural Language Supervision. In *International Conference on Machine Learning (ICML)*, pages 8748–8763, 2021.
- [40] Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. Zero-Shot Text-to-Image Generation. In *International Conference on Machine Learning (ICML)*, pages 8821–8831, 2021.
- [41] Ali Razavi, Aäron van den Oord, and Oriol Vinyals. Generating Diverse High-Fidelity Images with VQ-VAE-2. In *Conference on Neural Information Processing Systems (NeurIPS)*, pages 14837–14847, 2019.
- [42] Ariel Shamir and Ari Rappoport. Compacting oriental fonts by optimizing parametric elements. *The Visual Computer*, 15(6):302–318, 1999.
- [43] Yiren Song and Yuxuan Zhang. Clipfont: Text Guided Vector WordArt Generation. In *British Machine Vision Conference (BMVC)*, page 543, 2022.
- [44] Danyang Sun, Tongzheng Ren, Chongxuan Li, Hang Su, and Jun Zhu. Learning to Write Stylized Chinese Characters by Reading a Handful of Examples. In *International Joint Conference on Artificial Intelligence (IJCAI)*, volume abs/1712.06424, pages 920–927, 2018.
- [45] Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 2018.
- [46] Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jonathon Shlens, and Zbigniew Wojna. Rethinking the Inception Architecture for Computer Vision. In *Computer Vision and Pattern Recognition (CVPR)*, pages 2818–2826, 2016.

- [47] Licheng Tang, Yiyang Cai, Jiaming Liu, Zhibin Hong, Mingming Gong, Minhu Fan, Junyu Han, Jingtuo Liu, Errui Ding, and Jingdong Wang. Few-Shot Font Generation by Learning Fine-Grained Local Styles. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7895–7904, 2022.
- [48] Yuchen Tian. zi2zi: Master chinese calligraphy with conditional adversarial networks. <https://github.com/kaonashi-tyc/zi2zi>, 2017.
- [49] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is All you Need. In *Conference on Neural Information Processing Systems (NeurIPS)*, pages 5998–6008, 2017.
- [50] Chi Wang, Min Zhou, Tiezheng Ge, Yunling Jiang, Hujun Bao, and Weiwei Xu. Cf-Font: Content Fusion for Few-shot Font Generation. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [51] Yizhi Wang and Zhouhui Lian. Deepvecfont: synthesizing high-quality vector fonts via dual-modality learning. *ACM Transactions on Graphics (TOG)*, 40(6):265:1–265:15, 2021.
- [52] Z. Wang, A.C. Bovik, H.R. Sheikh, and E.P. Simoncelli. Image Quality Assessment: From Error Visibility to Structural Similarity. *IEEE Transactions on Image Processing*, 13(4):600–612, 2004.
- [53] S. Wu, Chih-Yuan Yang, and Jane Yung-jen Hsu. Calligan: Style and Structure-aware Chinese Calligraphy Character Generator. In *AI for Content Creation Workshop at CVPR (AI4CC)*, volume abs/2005.12500, 2020.
- [54] Yangchen Xie, Xinyuan Chen, Li Sun, and Yue Lu. Dg-Font: Deformable Generative Networks for Unsupervised Font Generation. In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 5126–5136, 2021.
- [55] Zhenhua Yang, Dezhi Peng, Yuxin Kong, Yuyi Zhang, Cong Yao, and Lianwen Jin. Fontdiffuser: One-Shot Font Generation via Denoising Diffusion with Multi-Scale Content Aggregation and Style Contrastive Learning. In *AAAI Conference on Artificial Intelligence (AAAI)*, volume abs/2312.12142, pages 6603–6611, 2024.
- [56] Lijun Yu, Yong Cheng, Zhiruo Wang, Vivek Kumar, Wolfgang Macherey, Yanping Huang, David A. Ross, Irfan Essa, Yonatan Bisk, Ming Yang, K. Murphy, A. Hauptmann, and Lu Jiang. Spae: Semantic Pyramid AutoEncoder for Multimodal Generation with Frozen LLMs. In *Thirty-seventh Conference on Neural Information Processing Systems*, volume abs/2306.17842, 2023.
- [57] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The Unreasonable Effectiveness of Deep Features as a Perceptual Metric. In *Computer Vision and Pattern Recognition (CVPR)*, pages 586–595, 2018.
- [58] Yexun Zhang, Ya Zhang, and Wenbin Cai. Separating Style and Content for Generalized Style Transfer. In *Computer Vision and Pattern Recognition (CVPR)*, pages 8447–8455. IEEE, 2018.
- [59] Yufan Zhang, Junkai Man, and Peng Sun. Mf-Net - A Novel Few-shot Stylized Multilingual Font Generation Method. In *Proceedings of the 30th ACM International Conference on Multimedia*. ACM, 2022.
- [60] Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A. Efros. Unpaired Image-to-Image Translation Using Cycle-Consistent Adversarial Networks. In *2017 IEEE International Conference on Computer Vision (ICCV)*, pages 2242–2251. IEEE, 2017.
- [61] Lei Zhu, Fangyun Wei, and Yanye Lu. Beyond Text: Frozen Large Language Models in Visual Signal Comprehension. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.

A Experiment Details

A.1 Implementation Details

IF-Font was trained in a server equipped with an Intel Xeon Silver 4110 CPU, 128 GB of RAM, and an NVIDIA Tesla V100 PCIe 16GB GPU. The training takes about 42 hours for 15 epochs with batch size of 128.

The VQGAN model in IF-Font loads pre-trained weights. It has a codebook size of 256, and the encoder downsamples the image with a factor of 8. Consequently, the length of the codebook indices corresponding to a single glyph is $l_T = 256$, whereas the IDS have a fixed length of $l_I = 35$. Our decoder consists of 10 Transformer blocks, each integrating a self-attention layer, a cross-attention layer, and a multi-layer perceptron (MLP). We have configured the model with 8 attention heads and a feature dimension of 384. In IF-Font, parameters are optimized using the AdamW optimizer [32], which employs a learning rate schedule that includes warmup and cosine annealing.

A.2 Training Data

We primarily obtain fonts from the [Foundertype](#) platform under a personal non-commercial academic research license. We present the examples of three IDS granularities in Table 5.

Table 5: Different IDS granularities. Mixed means that the IDSs of both granularities are used.

Granularity	IDS	Length
Componet	□✕ □✚田	5
Stroke	□□□□✓)✓ □□—□ □□□□ 7—□—	21
Mixed	□✕ □✚田[sep]□□□□✓)✓ □□—□ □□□□ 7—□—	27

B Additional Results

B.1 Further ablation of SSA

Table 6: Ablation studies on both branches of SSA.

SSA	FID \downarrow	L1 \downarrow	LPIPS \downarrow	RMSE \downarrow	SSIM \uparrow
w/o global	7.9766	0.1607	0.1347	0.3789	0.4775
w/o local	8.2578	0.1620	0.1364	0.3805	0.4756
full	8.4922	0.1597	0.1338	0.3775	0.4782

B.2 Radar plots

Fig. 10a depicts the one-shot setting, where style extraction poses a considerable challenge, leading to comparable performance across most models. Nevertheless, as the number of reference glyphs increases, the advantage of IF-Font becomes progressively apparent in Fig. 10b and 10c.

C Further discussion

C.1 Quantization accuracy

Since the target glyph is reconstructed from a quantized sequence, the accuracy of this reconstruction imposes a ceiling on IF-Font’s potential performance. As shown in Figure 9, there are slight differences between the target and the ground truth. Given that the glyphs are grayscale images and

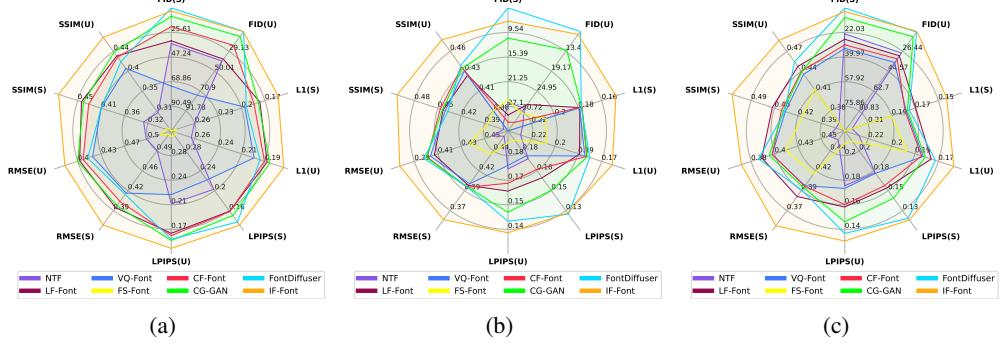


Figure 10: Compared with the methods based on the content-style disentanglement paradigm, IF-Font achieves state-of-the-art performance on all metrics under three few-shot settings. The metrics are annotated with brackets in the figure to specify the dataset used for evaluation: (S) represents UFSC, and (U) refers to UFUC. (a) 1-shot setting. (b) 3-shot setting. (c) 8-shot setting.

relatively simple, fine-tuning the VQ-GAN or switching to a superior quantized method is expected to minimize the loss of accuracy. We use original VQ-GAN checkpoints in order to highlight our main contributions.

C.2 The shortcomings of IDS

In fact, IDS is still not sufficiently perfect to identify Chinese characters. On the one hand, there are rule conflicts: a few Chinese characters are too similar, and their IDSs of stroke granularity are exactly the same. On the other hand, the spatial descriptions are insufficiently clear. For example, the left-right structure indicates that two components are placed on the left and right, but the distance between them is not specified, which requires the model to distinguish them through sufficient learning.

C.3 Complete results on novel fonts



Figure 11: Complete results generated by methods with novel fonts of Fig. 9.

We believe that generating novel fonts presents significant challenges for several reasons. First, they substantially deviate from standard character structures, involving a trade-off between the number of references and the quality of generation. Furthermore, font design involves subjectivity and randomness. For instance, in the first font depicted in Fig. 9, the position, size, and shape of the auspicious cloud patterns are the result of manual design. It is important to note that this issue is common across all font generation methods. Similar discussions can be found in LF-Font [38] and

CF-Font [50]. Unfortunately, this limitation has not been well addressed yet. We demonstrate the results of other methods on novel fonts in Fig. 11 to support the above conclusion.

C.4 Restrictions

To achieve high-quality generation, IF-Font requires training data to cover as many character components and font styles as possible. Due to the utilization of attention mechanisms, IF-Font entails quadratic computational complexity related to sequence length, while inference relies on autoregressive processes, resulting in a slow sampling speed.

D Details of the user study

To be accurate, we pay the Fuxi Youling Crowdsourcing Platform (a third-party platform) to conduct the user study. The participants in the user study are users of that platform, who will be informed of the full details in advance. We limit the number of participants to 30, and they are allowed to choose whether or not to participate and can complete the evaluation at any time and in any location using their own devices (usually a phone or computer).

All model names in the evaluation are replaced with numbers, and participants only need to select the option that best matches the ground truth from the model results displayed on the page. After the evaluation, the platform pays participants and provides us with de-identified model effectiveness data. We have no direct contact with participants and are unable to obtain their specific identities.

E Broader Impact

IF-Font could help to improve the productivity and creativity of font designers, and there is also hope for preserving ancient calligraphic works. In addition, IF-Font is capable of imitating the font style from as little as one reference glyph. That makes it easy to reproduce commercial font designs, raising concerns regarding potential copyright infringement. To safeguard the rights of font creators, we urge users to adhere to license requirements and call for the responsible use of our method.

F Licenses

We present a complete list of references and licenses in Table 7 for all the existing assets we used in this work.

Table 7: License information for the existing assets used.

Software Code	URL	License
VQGAN	Link	MIT license
nanoGPT	Link	MIT license
CG-GAN	Link	N/A
LF-Font	Link	MIT license
FS-Font	Link	MIT license
CF-Font	Link	N/A
VQ-Font	Link	N/A
NTF	Link	MIT license
torchmetrics	Link	Apache v2.0

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: The main claims made in the abstract and introduction accurately reflect the paper's contributions and scope.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We discuss the limitations in Section 5.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [NA]

Justification: This paper does not include theoretical results.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Please refer to Section 3 for detailed information on the model.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Code is available at <https://github.com/Stareven233/IF-Font>. We provide detailed instructions on how to obtain and process the data in Section 4.1 and Appendix A.2.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: We specify all the training and test details in Section 4.1 and Appendix A.1.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: Error bars are not reported because it would be too computationally expensive.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We provide the information on the computer resources in Appendix A.1.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Our research conform with the NeurIPS Code of Ethics in every respect.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader Impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: We discuss the societal impacts in Section 1 and Appendix E.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [Yes]

Justification: We describe safeguards in Appendix E.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Please refer to Appendix F for the references and licenses.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New Assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: This paper fully describes the details of assets in Section 3, 5 and Appendix A, F. The assets and documentation can be found in <https://github.com/Stareven233/IF-Font>.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and Research with Human Subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [Yes]

Justification: We include the details of the crowdsourcing experiment in Appendix D.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [Yes]

Justification: We include the details of the crowdsourcing experiment in Appendix D.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.